

Sender–Receiver Exercise 3: Reading for Senders

Harvard SEAS - Fall 2025

2025-10-09

1 Connected Components

The reading for receivers contains Section 1.1 below, as well as the statement of Theorem 2. Your goal will be to communicate the *proof* of Theorem 2 below to the receivers.

1.1 Connected Components

We begin by defining the *connected components* of an undirected graph. To gain intuition, you may find it useful to draw some pictures of graphs with multiple connected components and use them to help you follow along the proof.

Theorem 1. *Every undirected graph $G = (V, E)$ can be partitioned into connected components. That is, there are sets $V_0, \dots, V_{c-1} \subseteq V$ of vertices such that:*

1. V_0, \dots, V_{c-1} are disjoint, nonempty, and $V_0 \cup V_1 \cup \dots \cup V_{c-1} = V$. (This is what it means for V_0, \dots, V_{c-1} to be a partition of V .)
2. For every two vertices $u, v \in V$, u and v are in the same component V_i if and only if there is a path from u to v .

Moreover the sets V_0, \dots, V_{c-1} are unique (up to ordering), and are called the connected components of V .

In case you are interested, we include a proof of this theorem in the Hesterberg-Vadhan textbook (Section 12.5.3), but studying that proof is not required for this exercise.

We remark that for *directed* graphs, one can consider *weakly connected components*, where we ignore the directions of edges, and *strongly connected components*, where two vertices u, v are the same component if and only if there is a directed path from u to v and a directed path from v to u . Strongly connected components are more useful, but more complicated. In particular, unlike in undirected graphs (or weakly connected components), there can be edges crossing between strongly connected components.

1.2 Finding Connected Components via BFS

The main result of this exercise is an efficient algorithm for finding connected components:

Theorem 2. *There is an algorithm that, given an undirected graph $G = (V, E)$ with n vertices and m edges, partitions V into connected components in time $O(n + m)$.*

Proof. The idea is to do BFS from an arbitrary start vertex $s_0 \in V$, and let our first connected component V_0 consist of all the vertices that BFS finds. Then, if there are any vertices in $V \setminus V_0$, we pick an arbitrary $s_1 \in V \setminus V_0$, and do BFS from s_1 to identify a second component V_1 , and so on. Naively, this will give a runtime bound $O(c \cdot (n + m))$ where c is the number of connected components, because we do c executions of BFS, each of which could potentially take time $O(n + m)$.

We speed this up by showing that we can implement BFS from a start vertex s in time $O(n_s + m_s)$, where n_s and m_s are the number of vertices and edges, respectively, in the connected component containing s . Since we do BFS on distinct connected components (whose sets of vertices and edges are disjoint), our total runtime will just be $O(n + m)$.

However, it's not immediate that BFS from a start vertex s can be implemented in time $O(n_s + m_s)$.

Recall that our implementation of BFS maintained the set S of vertices already visited as an array of n bits (so that membership in S can be tested in constant time). If we re-initialize the array each time we run BFS, our run time will be at least n per BFS execution, regardless of how small the connected component containing s is.

Thus, we modify our description of BFS so that the bit-array S keeping track of the vertices we visit is already initialized as part of the input. It will also be convenient to allow us to use an arbitrary label ℓ to indicate which vertices we have visited in the current BFS execution rather than marking them with the bit 1; this will allow us to assign different labels for different connected components.

```

BFSlabel( $G, s, S, \ell$ ):
Input      : A directed graph  $G = (V, E)$ , a vertex  $s \in V$ , a label  $\ell \in \mathbb{N}$ , and an array
                $S$  of length  $n = |V|$  where for every vertex  $v$ ,  $S[v] \neq \ell$ 
Output    : The array is updated so that  $S[v] = \ell$  for every  $v$  reachable from  $s$ , and
               the other entries of  $S$  are unchanged
0  $S[s] = \ell$ ;
1  $F = \{s\}$  ;                               /* the frontier vertices */
2  $d = 0$ ;
3 /* loop invariant:  $S[v] = \ell$  iff  $v$  has distance  $\leq d$  from  $s$ ,  $F$  = vertices at
   distance  $d$  from  $s$  */
4 while  $F \neq \emptyset$  do
5    $F = \{v \in V : \exists u \in F \text{ s.t. } (u, v) \in E \text{ and } S[v] \neq \ell\}$  ;
6   foreach  $v \in F$  do  $S[v] = \ell$ ;
7    $d = d + 1$ ;

```

Algorithm .1: BFSlabel()

Similarly to the runtime analysis we did in the previous lecture, the runtime of BFSlabel(G, s, S, ℓ) can be bounded as

$$O\left(\sum_{d=0}^{\infty} \sum_{u \in F_d} (1 + \deg_{out}(u))\right) \leq O\left(\sum_{u \in R} (1 + \deg_{out}(u))\right).$$

where F_d is the set of vertices u such that $\text{dist}_G(s, u) = d$, and $R = \bigcup_{d=0}^{\infty} F_d$ is the set of vertices reachable from s .

Now the key point is that, in an undirected graph G , R is exactly the connected component containing s , so $|R| = n_s$ and $\sum_{u \in R} \deg_{out}(u) = 2m_s$ (since each of the m_s undirected edges

contributes to \deg_{out} for two vertices). Thus, the run time of $\text{BFSlabel}(G, s, S, \ell)$ is $O(n_s + m_s)$.

Now we can obtain our algorithm for connected components as follows:

BFS-CC(G):
Input : An undirected graph $G = (V, E)$
Output : The number ℓ of connected components in G and a partition of G into those components, specified by an array S of length $n = |V|$ with entries from $[\ell]$

```

0 Initialize  $S[v] = \star$  for all  $v \in V$ ;
1  $\ell = 0$ ;
2 foreach  $s \in V$  do
3   if  $S[s] = \star$  then
4     BFSlabel( $G, s, S, \ell$ );
5      $\ell = \ell + 1$ ;
6 return  $(\ell, S)$ 

```

Algorithm .2: BFS-CC()

To get some intuition for this algorithm, consider how it runs on some examples. For example, if we take the graph in the textbook's Shortest Walks Example 12.2 and Figure 1 (This is the same as figure 12.1 in the textbook).




Figure 1: BFS on the graph from Example 12.2 with $s = 0, t = 4$

and undirect all of its edges, one intermediate state of the array S would be $S = [0, 0, 0, 0, 0, 0, 1, \star, \star]$, when the algorithm has entirely explored one component and just started exploring the other. If we did the following swaps on vertex numbering $1 \leftrightarrow 7$ and $3 \leftrightarrow 8$ before running the algorithm, an intermediate state of the array would be $[0, 1, 0, \star, 0, 0, \star, 0, 0]$.

For the correctness of this algorithm, we prove the following loop invariant.

Claim .3. *At the start of each loop iteration, S has entries from $\{\star, 0, 1, \dots, \ell - 1\}$ with the vertices*

of each label $i \neq \star$ corresponding to a distinct connected component of G .

Proof of claim. We use induction on the number k of loop iterations that have been completed.

The base case ($k = 0$) follows because we initialize S to all \star 's.

For the induction step, assume that the claim is true at the start of a loop iteration k and we will argue that it is true at the start of loop iteration $k + 1$. The induction hypothesis tells us that at the start of loop iteration k , S has entries from $\{\star, 0, 1, \dots, \ell - 1\}$ with the vertices of each label $i \neq \star$ corresponding to a distinct connected component of G .

If $S[s] \neq \star$, then neither S nor ℓ change during loop iteration k , so the claim also holds at the start of loop iteration $k + 1$. If $S[s] = \star$, then S changes in Line 4 and we increment ℓ by 1 during loop iteration k . Since $S[s] = \star$, we know that s is not in any of the previously labelled connected components and thus $\text{BFSlabel}(G, s, S, \ell)$ will label the entire connected component of s with label ℓ and leave the rest of the array S unchanged. Thus, after Line 4, S has entries from $\{\star, 0, 1, \dots, \ell\}$ with the vertices of each label $i \neq \star$ corresponding to a distinct connected component of G . Since we increment ℓ , the claim will also hold at the start of loop iteration $k + 1$. \square

We also observe that due to the loop over $s \in V$, we will be sure to assign every vertex in V to some connected component.

For the runtime, observe all of the executions of Lines 0 to 3 take time $O(n)$ in total. Each time we run Line 4, we execute $\text{BFSlabel}(G, s, S, \ell)$, which runs in time $O(n_s + m_s)$, where n_s and m_s are the number of vertices and edges in the connected component of s . Since we run BFSlabel on vertices $s = s_0, \dots, s_{c-1}$ that are all in different connected components, the cost of all of these executions is

$$O\left(\sum_{i=0}^{c-1} (n_{s_i} + m_{s_i})\right) = O(n + m).$$

\square

In the above algorithm, BFS could have easily been replaced with another search strategy like depth-first search (DFS), since we don't care about finding *shortest* paths. It turns out that DFS can be used in a more sophisticated, two-pass fashion, to find the strongly connected components of a directed graph. That algorithm is covered in CS124.